


## Notebook for EDA and Dense Neural Network for Analysis of red wine quality

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten
```

### Loading data

```
from google.colab import drive
drive.mount('/content/drive')
```

```
file_path = '/content/drive/MyDrive/CS4824_Final/winequality-red.csv'
```

 Mounted at /content/drive

```
data = pd.read_csv(file_path)
```

## Data Dictionary

### **Response variable:** quality

- Description: Sensory score given by wine experts.
- Impact: Target variable; represents overall acceptance.

### **Explanatory variables**

#### fixed acidity:

- Description: Non-volatile acids that do not evaporate during fermentation.
- Impact: Contributes to the overall acidity and taste profile of the wine.

#### volatile acidity:

- Description: Volatile acids that can evaporate; primarily acetic acid.
- Impact: High levels may lead to an unpleasant vinegar taste.

#### citric acid:

- Description: A natural component of grapes; adds freshness and flavor.
- Impact: Enhances the sensory profile when in balanced amounts.

#### residual sugar:

- Description: Sugar remaining after fermentation; not converted to alcohol.
- Impact: Influences sweetness; higher levels make the wine sweeter.

#### chlorides:

- Description: Salt content in the wine.
- Impact: Excessive chlorides can result in a salty taste.

#### free sulfur dioxide:

- Description: SO<sub>2</sub> in the free form; acts as an antioxidant and antimicrobial agent.
- Impact: Protects wine from spoilage; too much can affect taste.

#### total sulfur dioxide:

- Description: Sum of free and bound sulfur dioxide.
- Impact: High levels can cause sensory defects and health concerns.

#### density:

- Description: Mass per unit volume; correlates with alcohol and sugar content.
- Impact: Used to monitor fermentation and quality control.

#### pH:

- Description: Scale of acidity (lower pH = more acidic).
- Impact: Affects color, stability, and taste.

sulphates :

- Description: Additive used to increase shelf life by preventing oxidation.
- Impact: Enhances flavor and longevity; excessive amounts can be detrimental.

alcohol :

- Description: Ethanol content in the wine.
- Impact: Affects the body, warmth, and caloric content.

## ✓ Work planned for Oct 21st to Oct 31st:

focus on **exploratory data analysis (EDA)** using Python visualization libraries such as Matplotlib and Seaborn to examine the dataset through histograms, scatter plots, box plots, pair plots, and heatmaps.

### Goal:

uncover key patterns, detect outliers, identify potential predictors, and gain initial insights into how the chemical properties of the wines relate to their quality ratings [link text](#).

We will proceed once variables fit the assumptions of our model.

## ✓ EDA - Process of examining distributions and preparing data for modeling

### Examining individual distributions

Head of dataframe

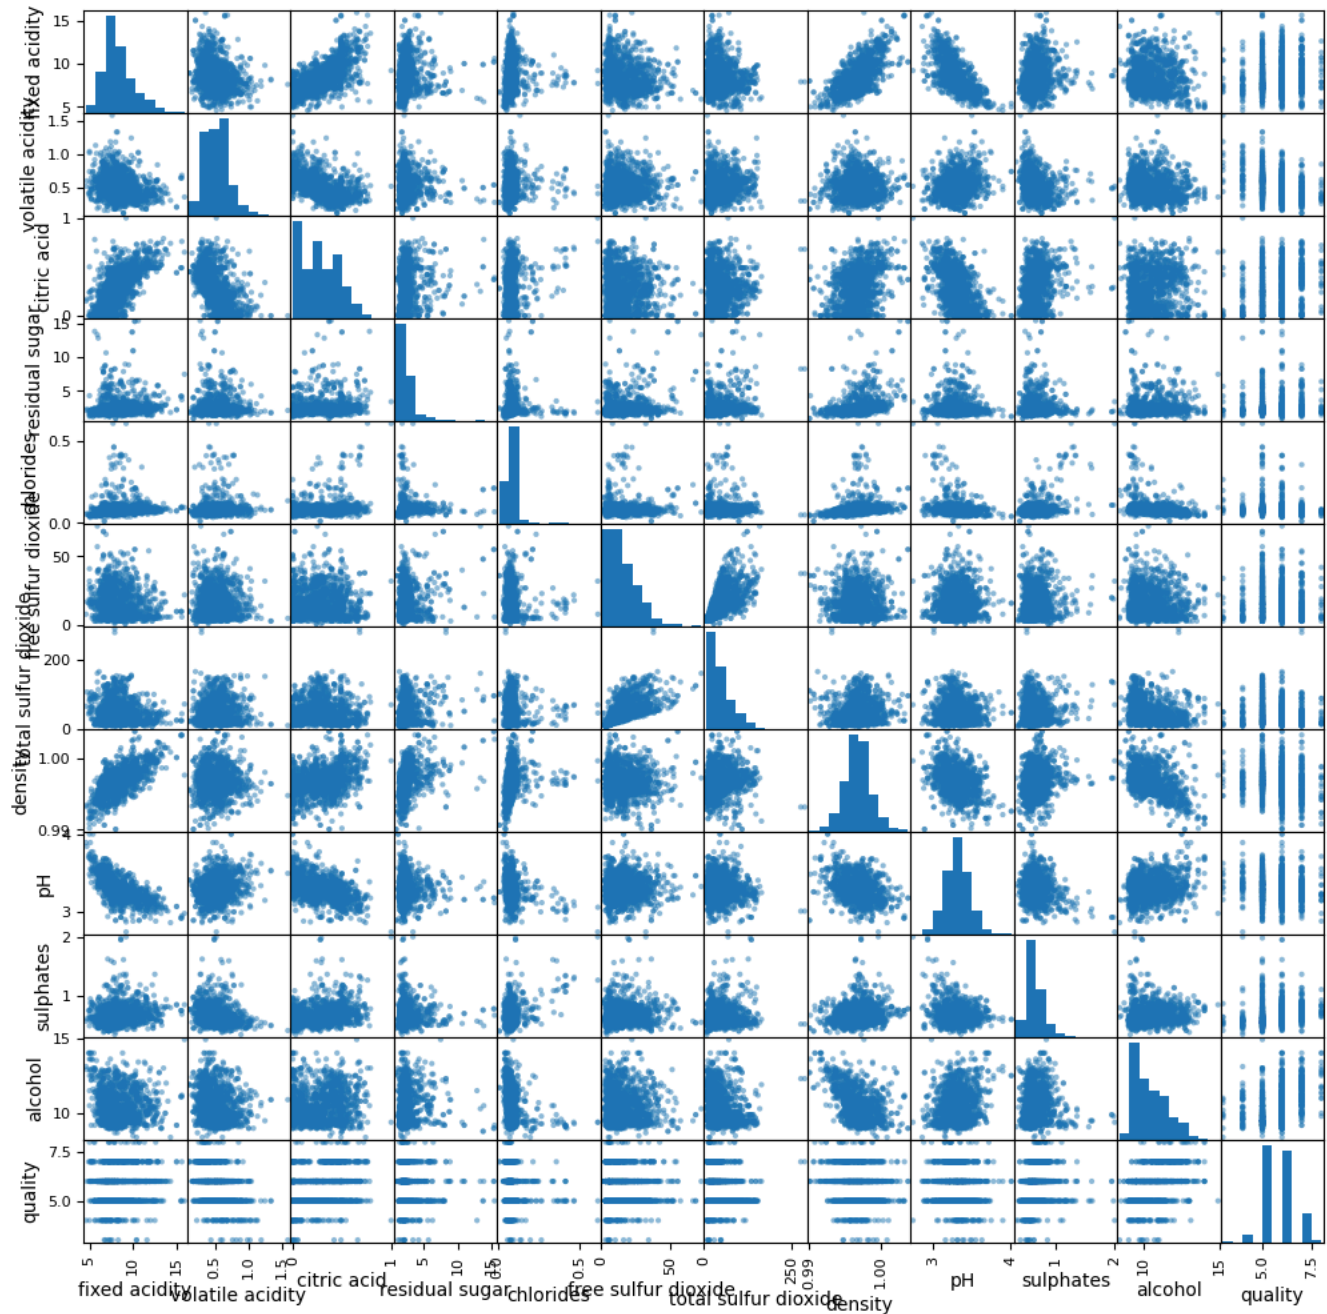
data[:20]

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
5	7.4	0.660	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4	5
6	7.9	0.600	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4	5
7	7.3	0.650	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0	7
8	7.8	0.580	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	9.5	7
9	7.5	0.500	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5	5
10	6.7	0.580	0.08	1.8	0.097	15.0	65.0	0.9959	3.28	0.54	9.2	5
11	7.5	0.500	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5	5
12	5.6	0.615	0.00	1.6	0.089	16.0	59.0	0.9943	3.58	0.52	9.9	5
13	7.8	0.610	0.29	1.6	0.114	9.0	29.0	0.9974	3.26	1.56	9.1	5
14	8.9	0.620	0.18	3.8	0.176	52.0	145.0	0.9986	3.16	0.88	9.2	5
15	8.9	0.620	0.19	3.9	0.170	51.0	148.0	0.9986	3.17	0.93	9.2	5
16	8.5	0.280	0.56	1.8	0.092	35.0	103.0	0.9969	3.30	0.75	10.5	7
17	8.1	0.560	0.28	1.7	0.368	16.0	56.0	0.9968	3.11	1.28	9.3	5
18	7.4	0.590	0.08	4.4	0.086	6.0	29.0	0.9974	3.38	0.50	9.0	4

## ▼ Scatterplot matrix for all variables

```
pd.plotting.scatter_matrix(data, figsize=(12,12))[-1]
```

```
array([[<Axes: xlabel='fixed acidity', ylabel='quality'>,
        <Axes: xlabel='volatile acidity', ylabel='quality'>,
        <Axes: xlabel='citric acid', ylabel='quality'>,
        <Axes: xlabel='residual sugar', ylabel='quality'>,
        <Axes: xlabel='chlorides', ylabel='quality'>,
        <Axes: xlabel='free sulfur dioxide', ylabel='quality'>,
        <Axes: xlabel='total sulfur dioxide', ylabel='quality'>,
        <Axes: xlabel='density', ylabel='quality'>,
        <Axes: xlabel='pH', ylabel='quality'>,
        <Axes: xlabel='sulphates', ylabel='quality'>,
        <Axes: xlabel='alcohol', ylabel='quality'>,
        <Axes: xlabel='quality', ylabel='quality'>], dtype=object)
```



## ▼ Individual distributions

We can use a normality test from the *sci.py* package to determine if any of our data are normally distributed

```

for item in data:
    print(str(item), stats.normaltest(data[item]))

fixed acidity NormaltestResult(statistic=224.53087840457746, pvalue=1.7528277735470436e-49)
volatile acidity NormaltestResult(statistic=143.41934355982863, pvalue=7.192589039756591e-32)
citric acid NormaltestResult(statistic=152.039214793795, pvalue=9.662822259281018e-34)
residual sugar NormaltestResult(statistic=1520.3239698236891, pvalue=0.0)
chlorides NormaltestResult(statistic=1783.1059225626427, pvalue=0.0)
free sulfur dioxide NormaltestResult(statistic=342.2591484251237, pvalue=4.779365332171615e-75)
total sulfur dioxide NormaltestResult(statistic=487.42725648953456, pvalue=1.4338908343436201e-106)
density NormaltestResult(statistic=30.70774994095191, pvalue=2.1473202738102222e-07)
pH NormaltestResult(statistic=33.684697471483915, pvalue=4.8468645347727716e-08)
sulphates NormaltestResult(statistic=906.8944479227036, pvalue=1.1759065222978855e-197)
alcohol NormaltestResult(statistic=154.17806951912513, pvalue=3.316328847318596e-34)
quality NormaltestResult(statistic=17.262400816355534, pvalue=0.00017845030333855057)

```

p-values are all very small, meaning that none of our variables are normally distributed, the variable closest to a normal distribution according to our test is quality, with  $p = 0.00018$

We can also use Q-Q plots to get a visual representation for normality


```

#draws a QQ-plot and performs shapiro normality test. Data is a list, title is a string
def test_normality(data, title):
    from scipy.stats import shapiro
    stats.probplot(x = data, dist='norm', plot = plt)
    plt.title('QQ Plot of ' + str(title))
    plt.show()
    stat, p = shapiro(data)
    print(f"Shapiro-Wilk Test Statistic for {title}: {stat}, p-value: {p}")


for item in data:
    if item != 'quality':
        test_normality(data[item], item)

```

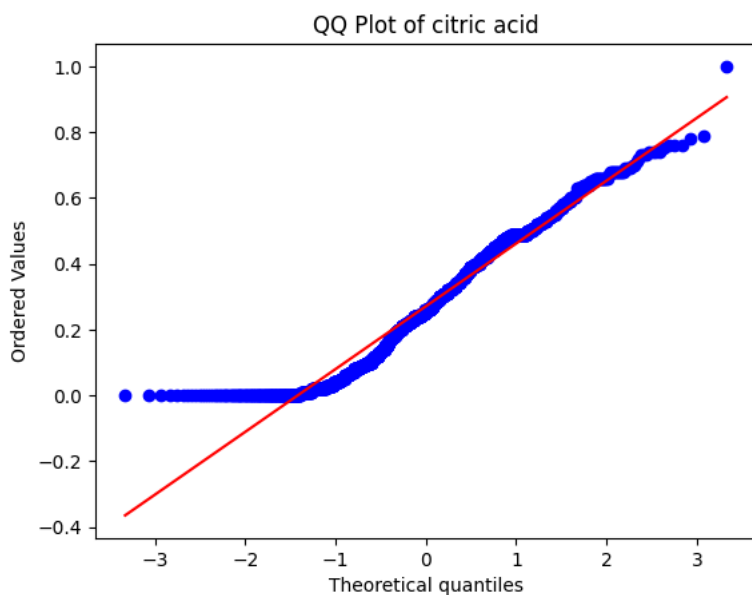
↓



Shapiro-Wilk Test Statistic for fixed acidity: 0.9420297903867135, p-value: 1.525011710791387e-24




Shapiro-Wilk Test Statistic for volatile acidity: 0.9743368805536368, p-value: 2.692934735712727e-16



Shapiro-Wilk Test Statistic for citric acid: 0.9552919890668837, p-value: 1.0219317829705018e-21


QQ Plot of residual sugar






Shapiro-Wilk Test Statistic for residual sugar: 0.5660771057163958, p-value: 1.0201616453237868e-52

QQ Plot of chlorides




Shapiro-Wilk Test Statistic for chlorides: 0.48424655122518334, p-value: 1.1790556953147118e-55


QQ Plot of free sulfur dioxide



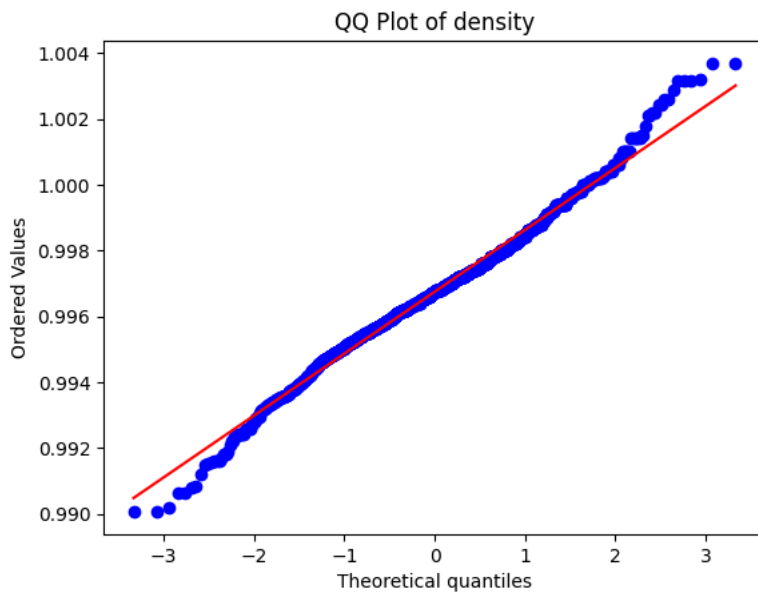
Shapiro-Wilk Test Statistic for free sulfur dioxide: 0.9018394916138583, p-value: 7.694596687816645e-31

QQ Plot of total sulfur dioxide







Shapiro-Wilk Test Statistic for total sulfur dioxide: 0.8732245604736051, p-value: 3.5734514102654424e-34




Shapiro-Wilk Test Statistic for density: 0.9908655166510911, p-value: 1.936052131352189e-08




Shapiro-Wilk Test Statistic for pH: 0.9934862934498192, p-value: 1.7122367757609613e-06





Shapiro-Wilk Test Statistic for sulphates: 0.8330437683911954, p-value: 5.823139712583187e-38



Shapiro-Wilk Test Statistic for alcohol: 0.9288390813054377, p-value: 6.644056905730039e-27



This gives us a better sense of how normally distributed our data is. It looks like most of the data are normal, and just deviate towards the tails. However, there are some variables, namely `residual sugar`, `chlorides`, `free sulfur dioxide`, `total sulfur dioxide`, `sulfates`, `alcohol` that should be modified to fit a normal distribution.

## ✓ Boxplots - examining outliers

We can create boxplots for each variable that effects our response to look for outliers, done so below


```
def draw_boxplot(df, item):
    sns.boxplot(x = df['quality'], y = df[item])
    plt.title('Boxplot grouped by quality for ' + str(item))
    plt.show()

def draw_scatterplot(df, item):
    sns.scatterplot(x = df['quality'], y = df[item])
    plt.title('Scatterplot grouped by quality for ' + str(item))
    plt.show()


for item in data:
    if item != 'quality':
        draw_boxplot(data, item)
        draw_scatterplot(data, item)
```

↓


Boxplot grouped by quality for fixed acidity




Scatterplot grouped by quality for fixed acidity


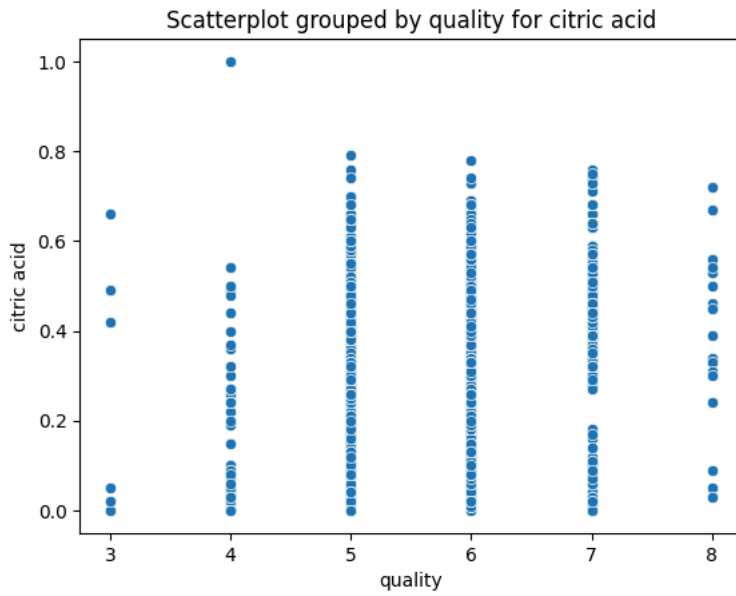
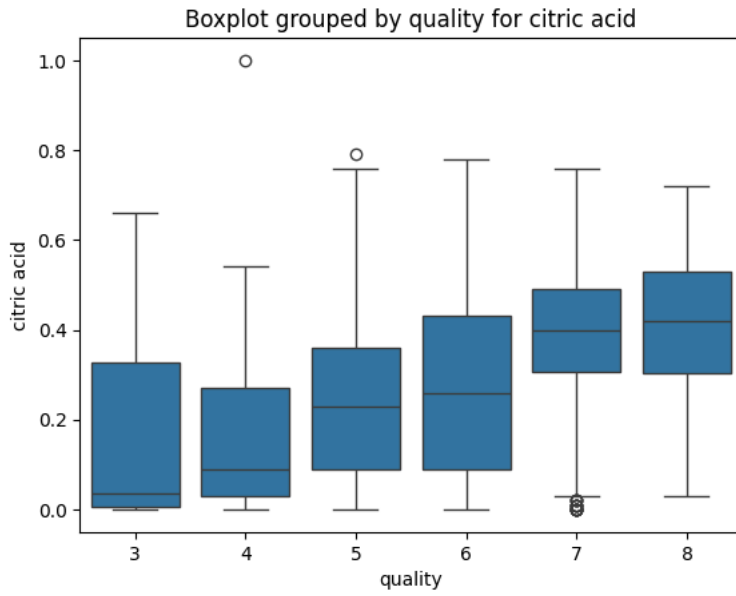



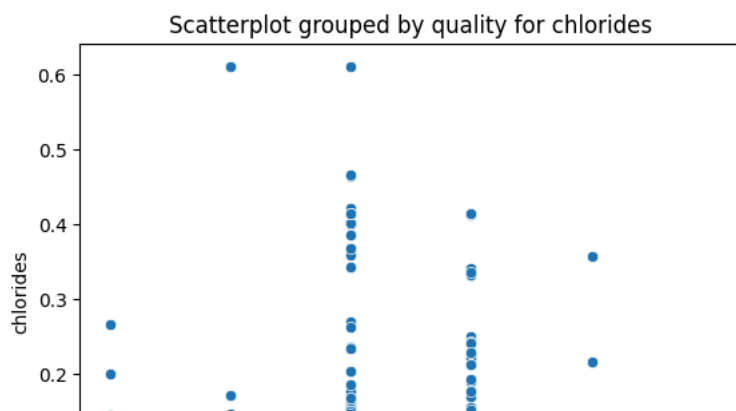
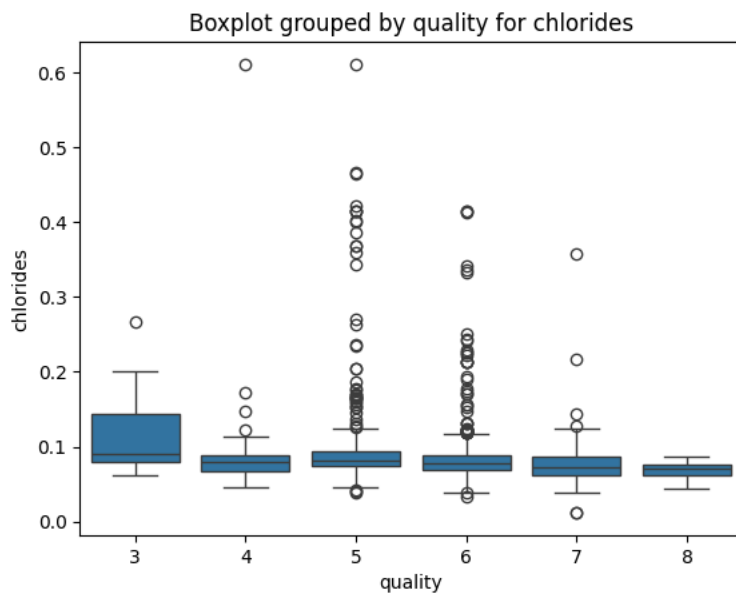
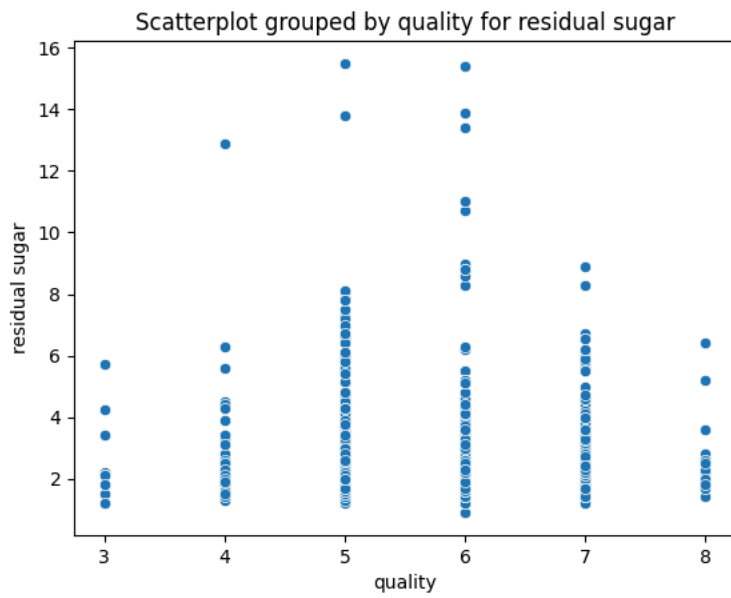

Boxplot grouped by quality for volatile acidity

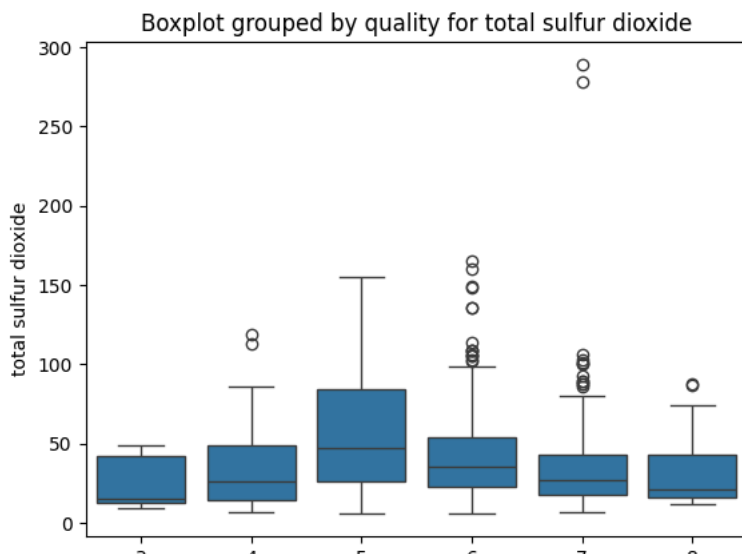
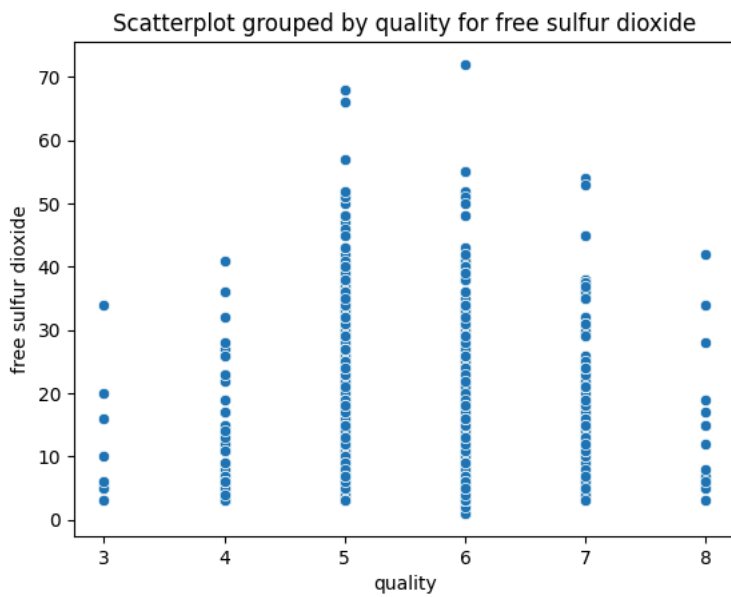
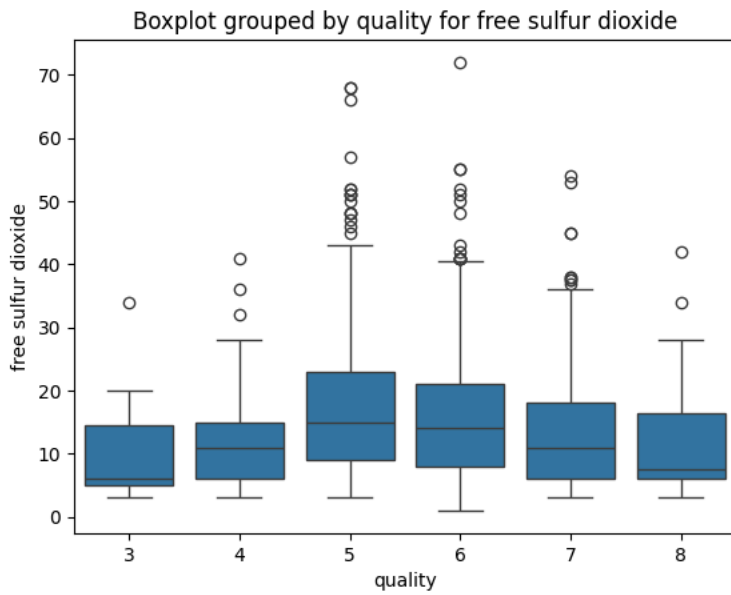







Scatterplot grouped by quality for volatile acidity

















Scatterplot grouped by quality for pH




Boxplot grouped by quality for sulphates




Scatterplot grouped by quality for sulphates






Boxplot grouped by quality for alcohol



Scatterplot grouped by quality for alcohol





## ✓ Data Transformation


Immediately, we can tell that we need to modify variables with a heavy right skew without looking at other visualizations. These variables are residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, and alcohol, which lines up pretty closely for what we determined should be changed when we created the Q-Q plots

### ✓ Residual Sugar

Below is what the distribution of residual sugar looks like, and a violin plot for how it relates to quality.

```
plt.hist(data['residual sugar'])
plt.show()

sns.violinplot(y = data['residual sugar'], x = 'quality', data = data)
plt.xlabel('Residual Sugar')
plt.ylabel('Quality')
```




Since our data is right-skew, a log-sqrt transformation will help us shift this to a normal distribution. We can visualize the change by plotting the Q-Q-Plots for the original and modified data, side-by-side, where we can see we were able to wrangle in some of the right-skew-ness


```
item = 'residual sugar'

vals = np.log(np.sqrt(data[item]))
```

```
test_normality(data[item], item)
test_normality(np.log(np.sqrt(data['residual sugar'])), 'log(residual sugar)')
```



Shapiro-Wilk Test Statistic for residual sugar: 0.5660771057163958, p-value: 1.0201616453237868e-52



Shapiro-Wilk Test Statistic for log(residual sugar): 0.8550735645340175, p-value: 5.4710630073218684e-36

because we exhibit a positive change for our data's distribution with this change, we will apply the following transformation to our code in-order to train our model.

Given  $\vec{S}$  is the residual sugar for all samples, we will take

$$\vec{S} \leftarrow \log(\vec{S})$$

```
residual_sugar_data = data['residual sugar']
data['residual sugar'] = np.log(residual_sugar_data)
```

## chlorides




Below is what the distribution of chlorides looks like, and a violin plot for how it relates to quality.

Double-click (or enter) to edit

```
plt.hist(data['chlorides'])  
plt.xlabel('Chlorides')  
plt.ylabel('Quality')  
plt.show()
```

```
sns.violinplot(y = data['chlorides'], x = 'quality', data = data)  
plt.xlabel('Chlorides')  
plt.ylabel('Quality')  
plt.show()
```

```
draw_boxplot(data, 'chlorides')
```



Fixed Acidity is relatively normal, with a slight right skew. Lets see how using log, sqrt, and boxcox (for observations with fixed acidity> 0) transformations change the shape of the QQ-plot and results of normality test.

```

from scipy.stats import boxcox


#draws a QQ-plot and performs shapiro normality test. Data is a list, title is a string
def test_normality(data, title):
    from scipy.stats import boxcox
    from scipy.stats import shapiro
    stats.probplot(x = data, dist='norm', plot = plt)
    plt.title('QQ Plot of ' + str(title))
    plt.show()
    stat, p = shapiro(data)
    print(f"Shapiro-Wilk Test Statistic for {title}: {stat}, p-value: {p}")

chlorides_boxcox, chlorides_lambda = boxcox(data['chlorides'])


test_normality(data['chlorides'], 'chlorides')
test_normality(np.log(data['chlorides']), 'log(chlorides)')
test_normality(np.sqrt(data['chlorides']), 'sqrt(chlorides)')
test_normality(chlorides_boxcox, 'boxcox(chlorides)')

```


[4]




Shapiro-Wilk Test Statistic for chlorides: 0.48424655122518334, p-value: 1.1790556953147118e-55




Shapiro-Wilk Test Statistic for log(chlorides): 0.8283642549328323, p-value: 2.3622721512330157e-38



Shapiro-Wilk Test Statistic for sqrt(chlorides): 0.6723604020744909, p-value: 4.071545937322437e-48





Shapiro-Wilk Test Statistic for boxcox(chlorides): 0.8716791986679291, p-value: 2.459048953286268e-34

Based on the results of our testing, it seems that boxcox did the best job of normalizing our data. We will change chlorides to boxcox(chlorides), which takes the following formula:

$$y_i^{(\lambda)} = \begin{cases} \frac{y_i^\lambda - 1}{\lambda}, & \text{if } \lambda \neq 0 \\ \ln(y_i), & \text{if } \lambda = 0 \end{cases}$$


We make this change in the next cell, along with storing the original value for `data['chlorides']` in `chlorides_data` for reference.

We can also observe two extreme outliers from the Q-Q-plot, with values around  $-14$ , while the rest of our values fall between  $(-8, 0)$ . We will also remove these.

```
chlorides_data = data['chlorides']
#data['chlorides'] = chlorides_boxcox
```

```
data[(data['chlorides'] < -10)]
```

```
draw_boxplot(data, 'chlorides')
```



```
data = data[(data['chlorides'] > -10)]
```

Finally, `chlorides` has the following shape, presented by a histogram, boxplot, and violin plot.





```
plt.hist(data['chlorides'])  
plt.xlabel('Chlorides')  
plt.ylabel('Quality')  
plt.show()
```


```
sns.violinplot(y = data['chlorides'], x = 'quality', data = data)  
plt.xlabel('Chlorides')  
plt.ylabel('Quality')  
plt.show()
```

```
draw_boxplot(data, 'chlorides')
```

```
test_normality(data['chlorides'], 'chlorides')
```







Shapiro-Wilk Test Statistic for chlorides: 0.48424655122518334, p-value: 1.1790556953147118e-55




## ▼ free sulfur dioxide

Below is the distribution of free sulfur dioxide, shown with a histogram, violin, and boxplot.

```
plt.hist(data['free sulfur dioxide'])
plt.xlabel('Free Sulfur Dioxide')
plt.ylabel('count')
plt.show()

sns.violinplot(y = data['free sulfur dioxide'], x = 'quality', data = data)
plt.ylabel('Free Sulfur Dioxide')
plt.xlabel('Quality')
plt.show()

draw_boxplot(data, 'free sulfur dioxide')
```




We have right skew data. We will look at the same transformations we made on chlorides to see if they help fix the skewness of our data. These are namely boxcox, sqrt, and log transformations.


```
free_sulfur_boxcox, free_sulfur_lambda = boxcox(data['free sulfur dioxide'])
```

```
test_normality(data['free sulfur dioxide'], 'free sulfur dioxide')  
test_normality(np.log(data['free sulfur dioxide']), 'log(free sulfur dioxide)')  
test_normality(np.sqrt(data['free sulfur dioxide']), 'sqrt(free sulfur dioxide)')  
test_normality(free_sulfur_boxcox, 'boxcox(free sulfur dioxide)')
```


[4]



Shapiro-Wilk Test Statistic for free sulfur dioxide: 0.9018394916138583, p-value: 7.694596687816645e-31




Shapiro-Wilk Test Statistic for log(free sulfur dioxide): 0.9833984364966966, p-value: 1.2463722155133395e-12



Shapiro-Wilk Test Statistic for sqrt(free sulfur dioxide): 0.9725749638438167, p-value: 6.699830655540897e-17

QQ Plot of boxcox(free sulfur dioxide)



Shapiro-Wilk Test Statistic for boxcox(free sulfur dioxide): 0.9874715597874553, p-value: 1.550940154035844e-10

Since both the log and boxcox transformations give a similar Shapiro test statistic ( $.98 < t < .99$ ), we will opt for a log transformation as it is slightly easier to interpret.

This means if  $\text{free sulfur dioxide} = \vec{S}$ , then our model with use


$$\vec{S} \leftarrow \log(\vec{S})$$

as the input for `free sulfur dioxide`. We make the change in the next cell

```
data['free sulfur dioxide'] = np.log(data['free sulfur dioxide'])
```

The last thing to do is re-assess for outliers, lets see how our normality and boxplots change when we remove the values that have `free sulfur dioxide = 0`.

```
test_normality(data[(data['free sulfur dioxide'] > 0.1)['free sulfur dioxide'], 'free sulfur dioxide'])
```



Shapiro-Wilk Test Statistic for free sulfur dioxide: 0.9822908798309025, p-value: 3.9812244852226943e-13




This actually makes our data less normal, so we will leave our data as is for now.

▼ total sulfur dioxide

```
plt.hist(data['total sulfur dioxide'])
plt.xlabel('Total Sulfur Dioxide')
plt.ylabel('Count')
plt.show()

sns.violinplot(y = data['total sulfur dioxide'], x = 'quality', data = data)
plt.ylabel('Total Sulfur Dioxide')
plt.xlabel('Quality')
plt.show()

draw_boxplot(data, 'total sulfur dioxide')
```



Before assessing normality, we can see that there are two extreme outliers for total sulfur dioxide. We will define these outliers as the cases where total sulfur dioxide is greater than 200. We will remove these.

```
data = data[(data['total sulfur dioxide']) < 200] #only include values < 200
```






Let's see how this has changed the plots from earlier, and if we should still try to normalize the skew of our data.

```
#same code from above
plt.hist(data['total sulfur dioxide'])
plt.xlabel('Total Sulfur Dioxide')
plt.ylabel('Count')
plt.show()

sns.violinplot(y = data['total sulfur dioxide'], x = 'quality', data = data)
plt.ylabel('Total Sulfur Dioxide')
plt.xlabel('Quality')
plt.show()

draw_boxplot(data, 'total sulfur dioxide')
```




Removing the outliers this way helped the skew become less drastic, but there is still an obvious right skew. We will look at boxcox, sqrt, and log transformations to see how they effect our normality.


```
total_sulfur_boxcox, total_sulfur_lambda = boxcox(data['total sulfur dioxide'])

test_normality(data['total sulfur dioxide'], 'total sulfur dioxide')
test_normality(np.log(data['total sulfur dioxide']**2), 'log(total sulfur dioxide)')
test_normality(np.sqrt(data['total sulfur dioxide']), 'sqrt(total sulfur dioxide)')
test_normality(total_sulfur_boxcox, 'boxcox(total sulfur dioxide)')
```


[4]



Shapiro-Wilk Test Statistic for total sulfur dioxide: 0.8901005442661531, p-value: 2.845716503915968e-32




Shapiro-Wilk Test Statistic for log(total sulfur dioxide): 0.987955192091651, p-value: 2.986940444373711e-10



Shapiro-Wilk Test Statistic for sqrt(total sulfur dioxide): 0.9643677622789827, p-value: 2.1870819867468594e-19

QQ Plot of boxcox(total sulfur dioxide)



Shapiro-Wilk Test Statistic for boxcox(total sulfur dioxide): 0.9885075208251186, p-value: 6.288487548570056e-10

We will choose to use a log transformation for our data, as it does nearly as well as boxcox while being more easily interpretable.

Given  $\vec{S}$  is the total sulfur dioxide for our samples, we will take

$$\vec{S} \leftarrow \log(\vec{S})$$

```
data['total sulfur dioxide'] = np.log(data['total sulfur dioxide'])
```




✓ alcohol

The last variable we will look at to change the distribution, based on our initial glance, will be alcohol.

```
#same code from above
plt.hist(data['alcohol'])
plt.xlabel('Alcohol (%)')
plt.ylabel('Count')
plt.show()

sns.violinplot(y = data['alcohol'], x = 'quality', data = data)
plt.ylabel('Alcohol (%)')
plt.xlabel('Quality')
plt.show()

draw_boxplot(data, 'alcohol')
```



Check how boxcox, log, and sqrt transformations effect outliers and normality.

```
alcohol_boxcox, alcohol_lambda = boxcox(data['alcohol'])
```


```
test_normality(data['alcohol'], 'alcohol')

test_normality(np.log(data['alcohol']), 'log(alcohol)')


test_normality(np.sqrt(data['alcohol']), 'sqrt(alcohol)')

test_normality(total_sulfur_boxcox, 'boxcox(alcohol)')
```


[4]



Shapiro-Wilk Test Statistic for alcohol: 0.9286791253859746, p-value: 6.482928724457048e-27




Shapiro-Wilk Test Statistic for log(alcohol): 0.9463831504798743, p-value: 1.1657602285248229e-23



Shapiro-Wilk Test Statistic for sqrt(alcohol): 0.938173968431395, p-value: 2.946880346198002e-25

QQ Plot of boxcox(alcohol)





Shapiro-Wilk Test Statistic for boxcox(alcohol): 0.9885075208251186, p-value: 6.288487548570056e-10

Based on the results of our testing, it seems that boxcox did the best job of normalizing our data. We will change `alcohol` to `boxcox(alcohol)`, which takes the following formula:

$$y_i^{(\lambda)} = \begin{cases} \frac{y_i^\lambda - 1}{\lambda}, & \text{if } \lambda \neq 0 \\ \ln(y_i), & \text{if } \lambda = 0 \end{cases}$$

We make this change in the next cell, along with storing the original value for `data['alcohol']` in `alcohol_data` for reference.

```
alcohol_data = data['alcohol']
data['alcohol'] = alcohol_boxcox
```




## ✎ sulfates

lets look at the shape of sulfates to see what changes should be made.

```
#same code from above
plt.hist(data['sulphates'])
plt.xlabel('Sulphates')
plt.ylabel('Count')
plt.show()

sns.violinplot(y = data['sulphates'], x = 'quality', data = data)
plt.ylabel('Sulphates')
plt.xlabel('Quality')
plt.show()

draw_boxplot(data, 'sulphates')
```




We will try our typical transformations, along with an inverse transformation to try and normalize the distribution of our data.

```
def test_varied_normality(data, title): #uses test_normality with boxcox, log, and sqrt transformations and the base case
    test_normality(data, title)
    test_normality(np.log(data), 'log(' + title + ')')
```


```
test_normality(np.sqrt(data), 'sqrt(' + title + ')')
test_normality(boxcox(data)[0], 'boxcox(' + title + ')')
```

```
test_varied_normality(data['sulphates'], 'sulphates')
test_normality(1/data['sulphates'], '1/sulphates')
```


[4]



Shapiro-Wilk Test Statistic for sulphates: 0.8331331995132054, p-value: 6.198155100264859e-38




Shapiro-Wilk Test Statistic for log(sulphates): 0.9589699699532543, p-value: 8.247531812929651e-21




Shapiro-Wilk Test Statistic for sqrt(sulphates): 0.9087192223058499, p-value: 6.562777046500386e-30

QQ Plot of boxcox(sulphates)



Shapiro-Wilk Test Statistic for boxcox(sulphates): 0.9964508719419674, p-value: 0.0009392978353782306

QQ Plot of 1/sulphates



Shapiro-Wilk Test Statistic for 1/sulphates: 0.9963963086800092, p-value: 0.0008240181117632388

Boxcox and Inverse perform similarly, so we will opt for the inverse transformation. Where if  $\vec{S}$  is sulphates, then our updated value becomes

$$\vec{S} = \frac{1}{S_i}$$

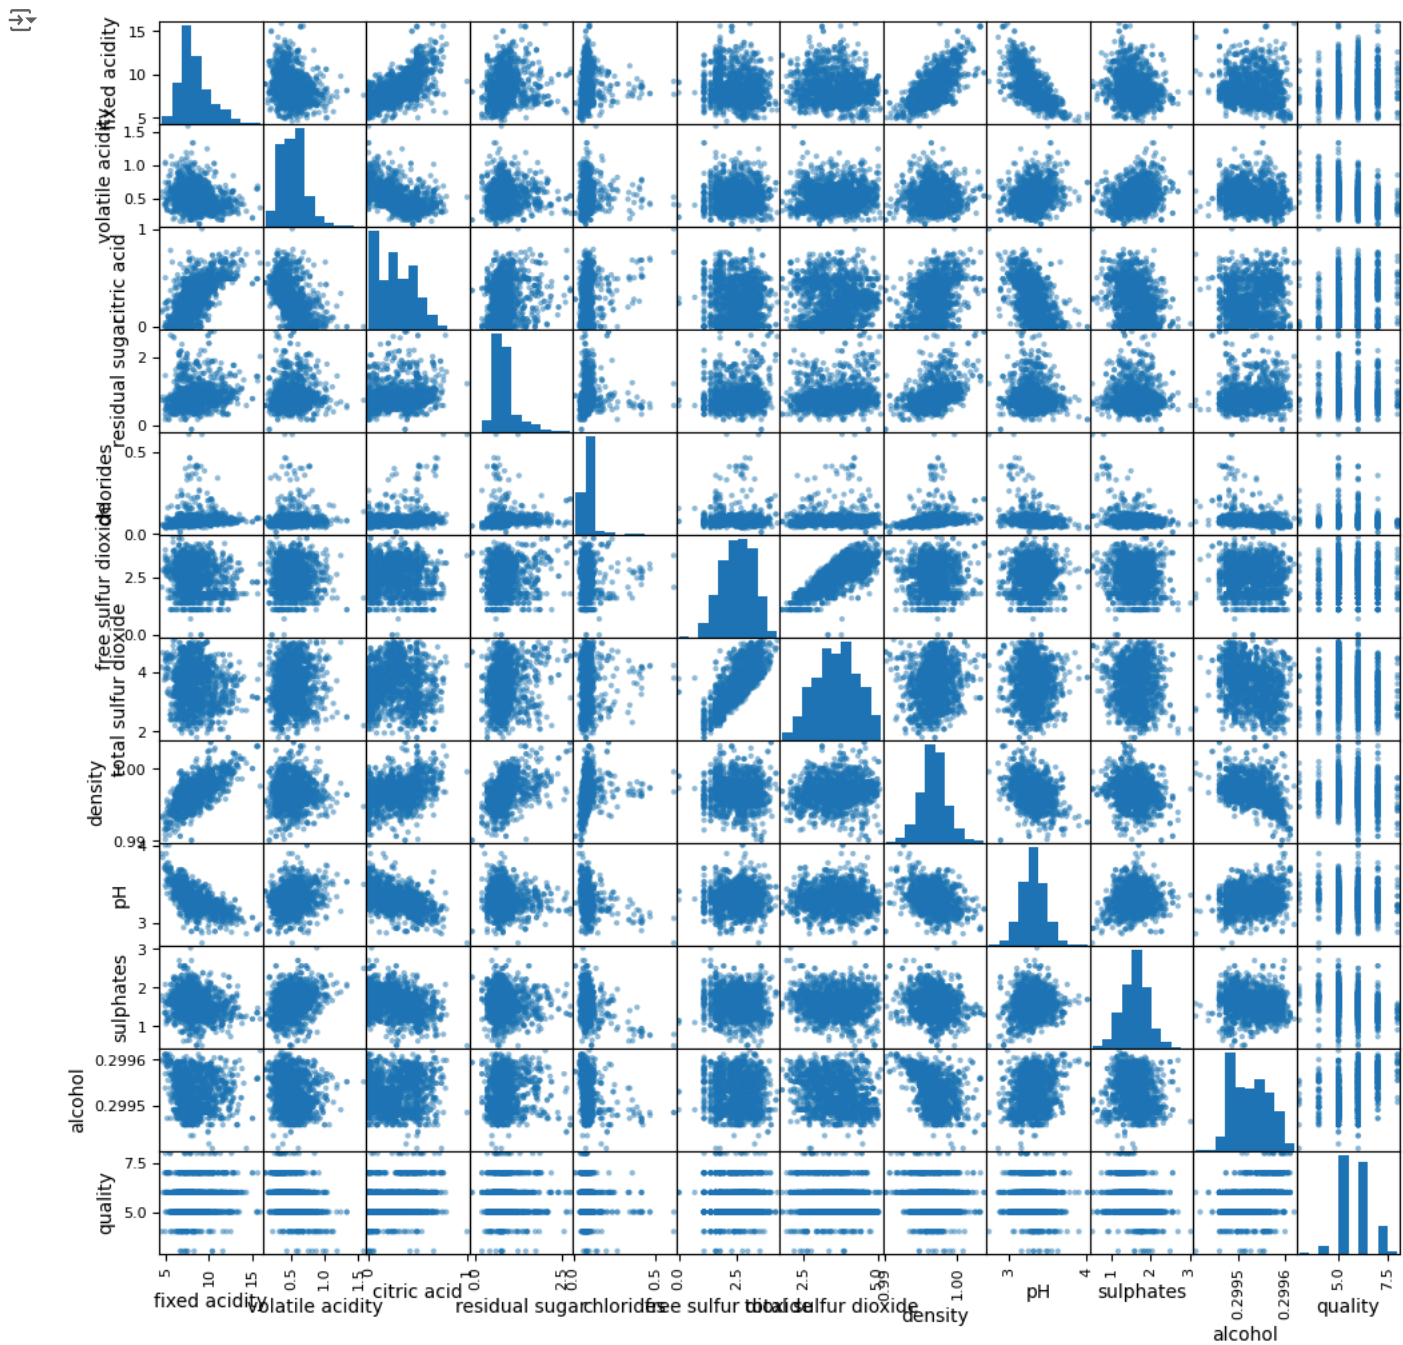
```
data['sulphates'] = 1/data['sulphates']
```

## ✓ Verifying data transformation

lets re-assess the scatter matrix and see if we have done enough with modifying distributions

## ✓ Scatterplot Matrix

```
pd.plotting.scatter_matrix(data, figsize=(12,12))[-1]
plt.show()
```







## ✓ Boxplots for Outliers





Double-click (or enter) to edit





```
for item in data:  
    if item != 'quality':  
        draw_boxplot(data, item)
```



(4)











Based on the boxplots, there are only two glaring outliers. There is one case where volatile acidity is  $\sim 1.6$ , where the rest of the values fall between  $(0.1, 1.4)$ . Similarly, citric acid falls between  $(0, 0.8)$ , and there is an observation with citric acid = 1. These are the last changes that will be made before we apply standard scaler.

```
data = data[(data['volatile acidity'] < 1.5)] # removing outliers for volatile acidity
data = data[(data['citric acid'] < 0.95)]
```

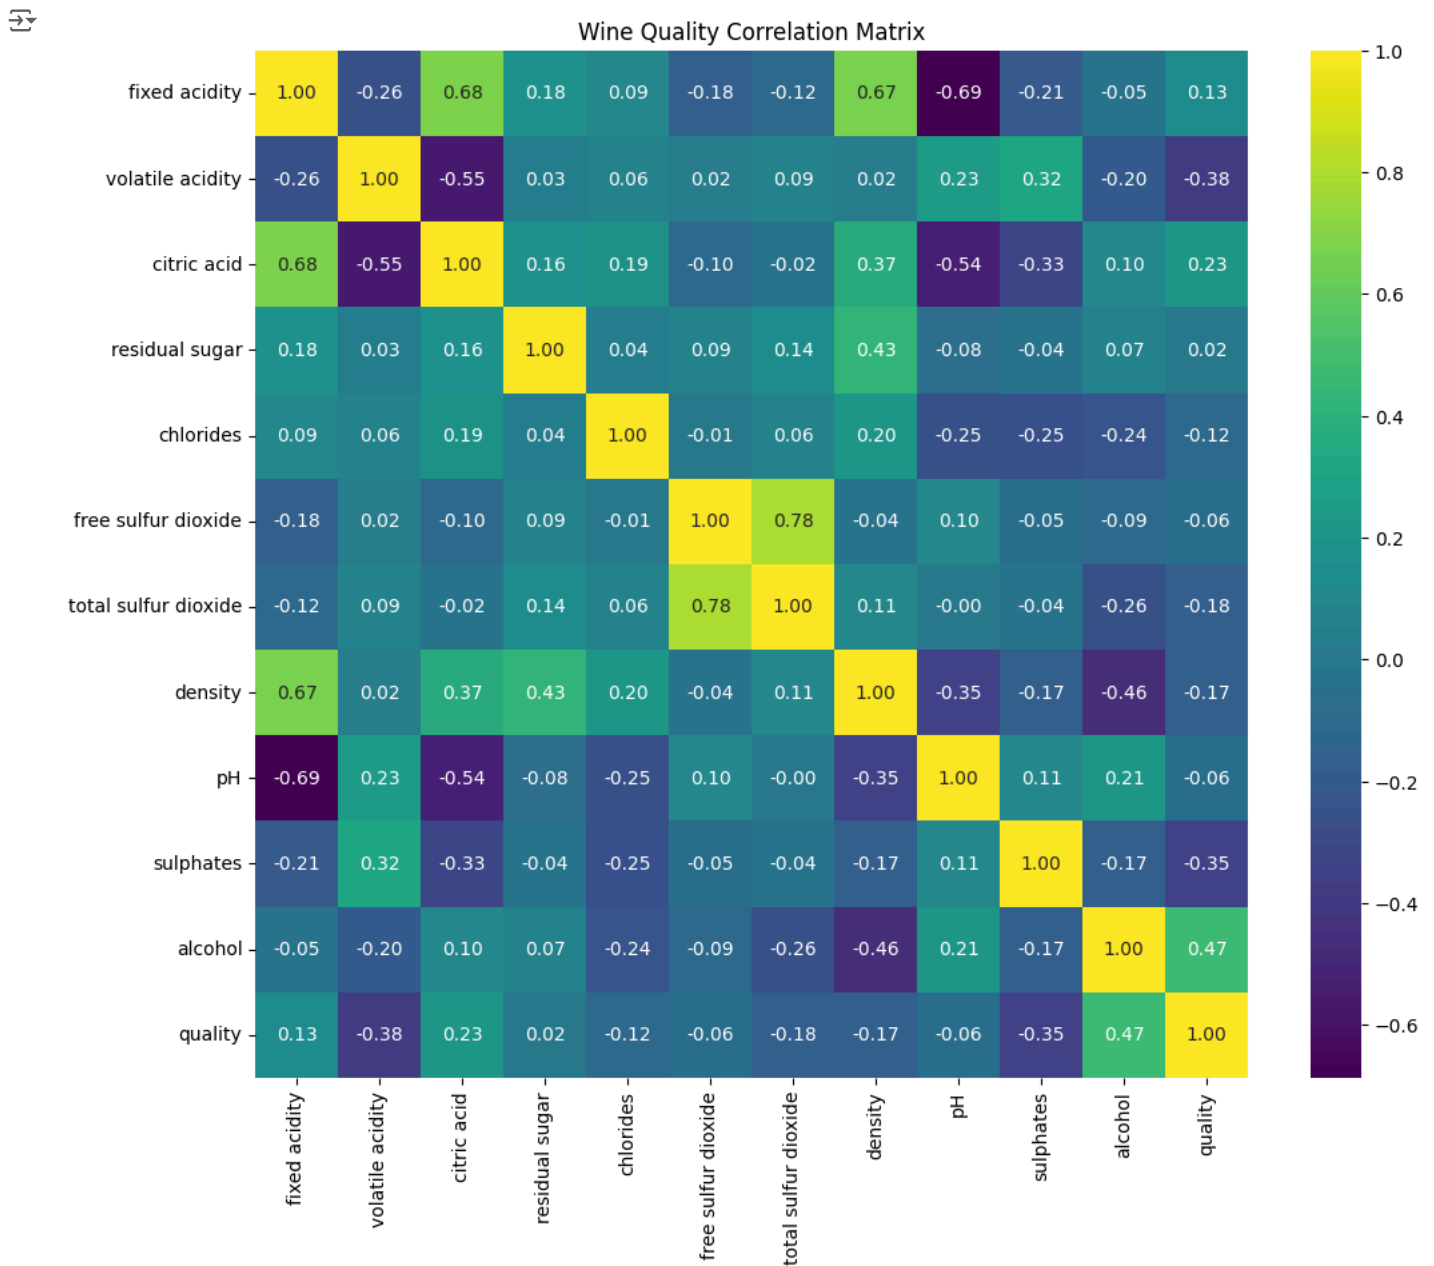
## ✓ Work Planned for November 1st to November 7th

Project will proceed with data processing. We will **apply the StandardScaler** function from TensorFlow to normalize the features and carry out feature engineering to create new variables that may enhance the model's ability to predict wine quality. This step is *essential for preparing the data* for neural network training by ensuring all features are appropriately scaled and engineered.

## ✓ Correlation Matrix

It's important to check the covariance matrix for confounding variables. If two vars are highly correlated, it might be wise to only include one of them in the model, or create a new feature all together.

```
plt.figure(figsize=(12, 10))
sns.heatmap(data.corr(), annot=True, fmt=".2f", cmap='viridis')
plt.title('Wine Quality Correlation Matrix')
plt.show()
```



Since  $corr(\text{free sulfur dioxide}, \text{total sulfur dioxide}) = 0.78$ , we will remove free sulfur dioxide from the dataset for training the model, since it will be partly included with total sulfur dioxide.

## ✓ Train-Test split

```
from sklearn.model_selection import train_test_split

X = data.drop('quality', axis=1)
y = data['quality']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
```

Below is duplicated code with unique naming, that will be used at the end for validation

```
X = data.drop('quality', axis=1)
y = data['quality']
X_training, X_testing, y_training, y_testing = train_test_split(X, y, test_size=0.2, random_state=1)
```

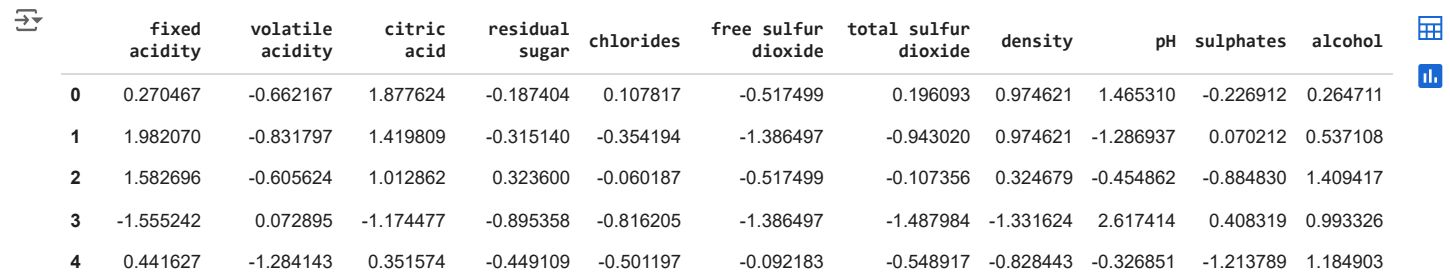
## ✓ Applying StandardScaler

```

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
# Transform test data
X_test_scaled = scaler.transform(X_test)
X_test_scaled = pd.DataFrame(X_test_scaled, columns=X_test.columns)
X_test_scaled.head()

```



	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	0.270467	-0.662167	1.877624	-0.187404	0.107817	-0.517499	0.196093	0.974621	1.465310	-0.226912	0.264711
1	1.982070	-0.831797	1.419809	-0.315140	-0.354194	-1.386497	-0.943020	0.974621	-1.286937	0.070212	0.537108
2	1.582696	-0.605624	1.012862	0.323600	-0.060187	-0.517499	-0.107356	0.324679	-0.454862	-0.884830	1.409417
3	-1.555242	0.072895	-1.174477	-0.895358	-0.816205	-1.386497	-1.487984	-1.331624	2.617414	0.408319	0.993326
4	0.441627	-1.284143	0.351574	-0.449109	-0.501197	-0.092183	-0.548917	-0.828443	-0.326851	-1.213789	1.184903

Next steps: [Generate code with X\\_test\\_scaled](#) [View recommended plots](#) [New interactive sheet](#)


## Updated Covariance Matrix

Since our values are now scaled, our covariance matrix will tell us more about the data and look much closer to the correlation matrix on the non-scaled data

```

plt.figure(figsize=(12, 10))
sns.heatmap(X_test_scaled.cov(), annot=True, fmt=".2f", cmap='viridis')
plt.title('Wine Quality Correlation Matrix')
plt.show()

```



## Random Forest for feature importance


```
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier
model = RandomForestRegressor(n_estimators=100, random_state=1)
# If classification, use:
# model = RandomForestClassifier(n_estimators=100, random_state=42)

# Fit the model
model.fit(X, y)

# Get feature importances
importances = model.feature_importances_
feature_names = X.columns
feature_importance_df = pd.DataFrame({
    'feature': feature_names,
    'importance': importances
}).sort_values('importance', ascending=False)

print(feature_importance_df)
plt.figure(figsize=(10, 6))
sns.barplot(x='importance', y='feature', data=feature_importance_df)
plt.title('Random Forest Feature Importance')
plt.show()
```

	feature	importance
10	alcohol	0.279491
9	sulphates	0.142195
1	volatile acidity	0.118382
6	total sulfur dioxide	0.081100
4	chlorides	0.064307
8	pH	0.059135
3	residual sugar	0.056924
7	density	0.053345
5	free sulfur dioxide	0.049042
0	fixed acidity	0.048168
2	citric acid	0.047911



All of our features seem to be at least somewhat important to predicting `quality`, so we will leave our features as is for now, and focus on feature engineering more later on.

## ✓ Work Planned for November 8th to November 15th

The project focus will shift to model selection and preliminary testing, where we will develop an initial deep neural network model using TensorFlow's Keras API. This stage will involve defining the network architecture, selecting activation functions, and running early tests to evaluate the model's initial performance on the dataset. We will validate the dataset for compatibility with the neural network and assess any core issues or necessary adjustments.

We also need to perform one-hot encoding on our `quality` classes, since they go from 3 to 8, instead of starting at zero.

## ✓ Initial Model selection, compilation, training, and sharing accuracy

We will begin with a simple DNN model trained on one layer, and potentially add as we go. We will use ReLU activation at first, but we will experiment with different activation functions.

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

#n_classes = y_train.shape[1] #currently not used
n_features = X_train.shape[1]

#initializing the model design
model = keras.Sequential([
    layers.Input(shape=(n_features,)),
```



```

# Hidden Layer
layers.Dense(64, activation='relu'),
layers.Dense(1) # Output layer for regression (linear by default), n_classes, activation = 'softmax' for classification
])

'''
Compile the model for training. Since this is a regression
problem, we use the 'mse' loss function and 'mae' as
the desired performance metric.
'''
opt = 'adam'
model.compile(loss='mse',
              optimizer=opt,
              metrics=['mae'])

```

## ▼ Initial Model Training

Before training the model, we need to partition our training data in validation and training.

```

# partition training set into training and validation set
X_validate = X_train[1000:]
X_train = X_train[:1000]
y_validate = y_train[1000:]
y_train = y_train[:1000]

print(f"X_train has shape {X_train.shape}")
print(f"y_train has shape {y_train.shape}")
print(f"X_validate has shape {X_validate.shape}")
print(f"y_validate has shape {y_validate.shape}")

```

```

↔ X_train has shape (1000, 11)
   y_train has shape (1000,)
   X_validate has shape (276, 11)
   y_validate has shape (276,)

```

#change the values from going

```

history_init_model = model.fit(
    X_train, y_train,
    epochs=35,
    validation_data=(X_validate, y_validate),
    batch_size=32,
    verbose=1
)

# Plot training accuracy
plt.plot(history_init_model.history['mae'], 'o-', label='DNN + {} activation'.format('relu'))
plt.title('training accuracy')
plt.ylabel('training accuracy')
plt.xlabel('epoch')
plt.legend(loc='upper right')
plt.show()

#plot validation accuracy
plt.plot(history_init_model.history['val_mae'], 'o-', label='DNN + {} activation'.format('relu'))
plt.title('validation accuracy')
plt.ylabel('validation accuracy')
plt.xlabel('epoch')
plt.legend(loc='upper right')
plt.show()

```

Epoch 1/35  
32/32 ————— 2s 33ms/step - loss: 12.8512 - mae: 3.3354 - val\_loss: 1.0198 - val\_mae: 0.8070

Epoch 2/35  
32/32 ————— 0s 3ms/step - loss: 1.1617 - mae: 0.8653 - val\_loss: 0.9545 - val\_mae: 0.7814

Epoch 3/35  
32/32 ————— 0s 2ms/step - loss: 1.0596 - mae: 0.8267 - val\_loss: 0.9000 - val\_mae: 0.7603

Epoch 4/35  
32/32 ————— 0s 4ms/step - loss: 0.9662 - mae: 0.7943 - val\_loss: 0.8797 - val\_mae: 0.7520

Epoch 5/35  
32/32 ————— 0s 3ms/step - loss: 0.9722 - mae: 0.8025 - val\_loss: 0.8793 - val\_mae: 0.7523

Epoch 6/35  
32/32 ————— 0s 2ms/step - loss: 0.9597 - mae: 0.7913 - val\_loss: 0.8427 - val\_mae: 0.7373

Epoch 7/35  
32/32 ————— 0s 3ms/step - loss: 0.9449 - mae: 0.7887 - val\_loss: 0.8222 - val\_mae: 0.7296

Epoch 8/35  
32/32 ————— 0s 3ms/step - loss: 0.8980 - mae: 0.7706 - val\_loss: 0.8055 - val\_mae: 0.7240

Epoch 9/35  
32/32 ————— 0s 3ms/step - loss: 0.8965 - mae: 0.7674 - val\_loss: 0.8081 - val\_mae: 0.7256

Epoch 10/35  
32/32 ————— 0s 4ms/step - loss: 0.9013 - mae: 0.7582 - val\_loss: 0.7707 - val\_mae: 0.7108

Epoch 11/35  
32/32 ————— 0s 3ms/step - loss: 0.7868 - mae: 0.7269 - val\_loss: 0.7543 - val\_mae: 0.7039

Epoch 12/35  
32/32 ————— 0s 3ms/step - loss: 0.7515 - mae: 0.7143 - val\_loss: 0.7402 - val\_mae: 0.6972

Epoch 13/35  
32/32 ————— 0s 4ms/step - loss: 0.8512 - mae: 0.7380 - val\_loss: 0.7296 - val\_mae: 0.6949

Epoch 14/35  
32/32 ————— 0s 3ms/step - loss: 0.7693 - mae: 0.7143 - val\_loss: 0.7290 - val\_mae: 0.6946

Epoch 15/35  
32/32 ————— 0s 3ms/step - loss: 0.7329 - mae: 0.7015 - val\_loss: 0.6901 - val\_mae: 0.6796

Epoch 16/35  
32/32 ————— 0s 3ms/step - loss: 0.7125 - mae: 0.6943 - val\_loss: 0.6720 - val\_mae: 0.6722

Epoch 17/35  
32/32 ————— 0s 3ms/step - loss: 0.7155 - mae: 0.6860 - val\_loss: 0.6562 - val\_mae: 0.6656

Epoch 18/35  
32/32 ————— 0s 3ms/step - loss: 0.7022 - mae: 0.6828 - val\_loss: 0.6532 - val\_mae: 0.6642

Epoch 19/35  
32/32 ————— 0s 3ms/step - loss: 0.6501 - mae: 0.6632 - val\_loss: 0.6284 - val\_mae: 0.6528

Epoch 20/35  
32/32 ————— 0s 3ms/step - loss: 0.6928 - mae: 0.6865 - val\_loss: 0.6341 - val\_mae: 0.6527

Epoch 21/35  
32/32 ————— 0s 4ms/step - loss: 0.6411 - mae: 0.6618 - val\_loss: 0.6112 - val\_mae: 0.6421

Epoch 22/35  
32/32 ————— 0s 5ms/step - loss: 0.5566 - mae: 0.6148 - val\_loss: 0.5933 - val\_mae: 0.6308

Epoch 23/35  
32/32 ————— 0s 2ms/step - loss: 0.5870 - mae: 0.6356 - val\_loss: 0.5916 - val\_mae: 0.6221

Epoch 24/35  
32/32 ————— 0s 3ms/step - loss: 0.6032 - mae: 0.6432 - val\_loss: 0.5806 - val\_mae: 0.6208

Epoch 25/35  
32/32 ————— 0s 4ms/step - loss: 0.5834 - mae: 0.6288 - val\_loss: 0.5706 - val\_mae: 0.6126

Epoch 26/35  
32/32 ————— 0s 3ms/step - loss: 0.6172 - mae: 0.6277 - val\_loss: 0.6486 - val\_mae: 0.6454

Epoch 27/35  
32/32 ————— 0s 5ms/step - loss: 0.5792 - mae: 0.6018 - val\_loss: 0.5595 - val\_mae: 0.6028

Epoch 28/35  
32/32 ————— 0s 4ms/step - loss: 0.5567 - mae: 0.6078 - val\_loss: 0.5530 - val\_mae: 0.5964

Epoch 29/35  
32/32 ————— 0s 3ms/step - loss: 0.5160 - mae: 0.5861 - val\_loss: 0.5760 - val\_mae: 0.6081

Epoch 30/35  
32/32 ————— 0s 4ms/step - loss: 0.5843 - mae: 0.6246 - val\_loss: 0.5572 - val\_mae: 0.5974

Epoch 31/35  
32/32 ————— 0s 3ms/step - loss: 0.5387 - mae: 0.5859 - val\_loss: 0.5465 - val\_mae: 0.5891

Epoch 32/35  
32/32 ————— 0s 3ms/step - loss: 0.6015 - mae: 0.6031 - val\_loss: 0.5609 - val\_mae: 0.5975

Epoch 33/35  
32/32 ————— 0s 3ms/step - loss: 0.5186 - mae: 0.5643 - val\_loss: 0.5613 - val\_mae: 0.5963

Epoch 34/35  
32/32 ————— 0s 4ms/step - loss: 0.5151 - mae: 0.5769 - val\_loss: 0.5553 - val\_mae: 0.5859

Epoch 35/35  
32/32 ————— 0s 4ms/step - loss: 0.5678 - mae: 0.5920 - val\_loss: 0.5480 - val\_mae: 0.5878

